

Minimum Spanning Tree

- Terminology and Properties
- Prim's Algorithm
- Union-Find
- Kruskal's Algorithm
- Baruvka's Algorithm
- Traveling Salesperson Problem

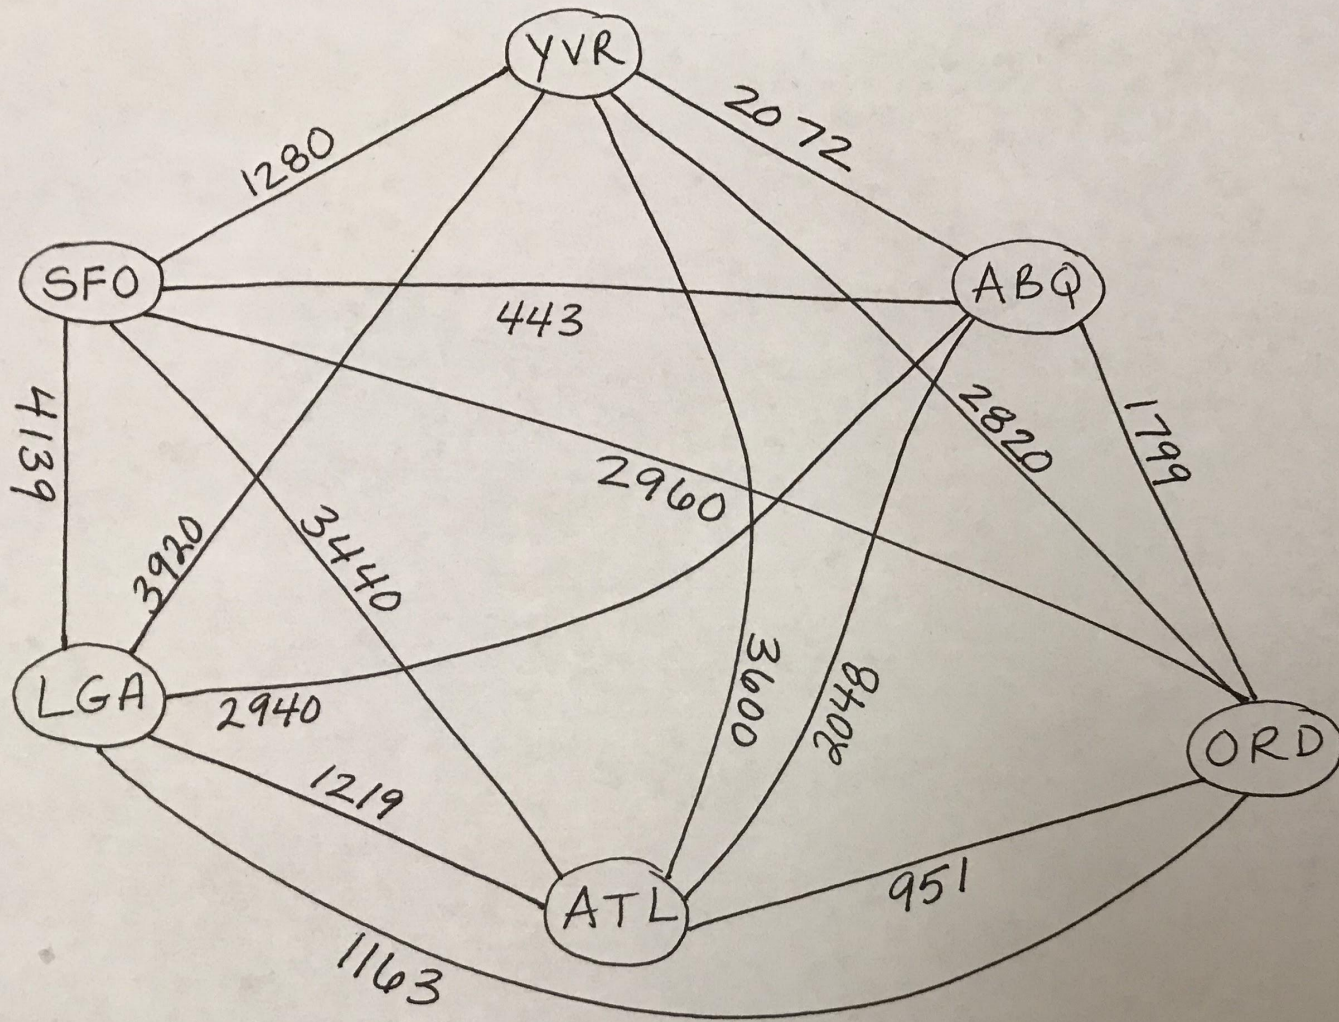
Imagine...

You are employed at a brand new airline, and your first task is to determine which flights the airline should provide according to the following guidelines:

- There needs to be a way to get between each pair of cities/airports in the following list: San Francisco, CA (SFO); Vancouver, BC (YVR); Albuquerque, NM (ABQ); Chicago, IL (ORD); Atlanta, GA (ATL); New York, NY (LGA)
- The route between two cities does not have to be a direct flight.
- The total number of flights should be minimized.
- The total distance covered by all flights should be minimized.
- Assume that if a flight exists, it goes both ways (e.g. SFO→YVR and YVR→SFO)

How would you go about solving this problem?

	SFO	YVR	ABQ	ORD	ATL	LGA
SFO	0	1280	443	2960	3440	4139
YVR	1280	0	2072	2820	3600	3920
ABQ	443	2072	0	1799	2048	2940
ORD	2960	2820	1799	0	951	1163
ATL	3440	3600	2048	951	0	1219
LGA	4139	3920	2940	1163	1219	0



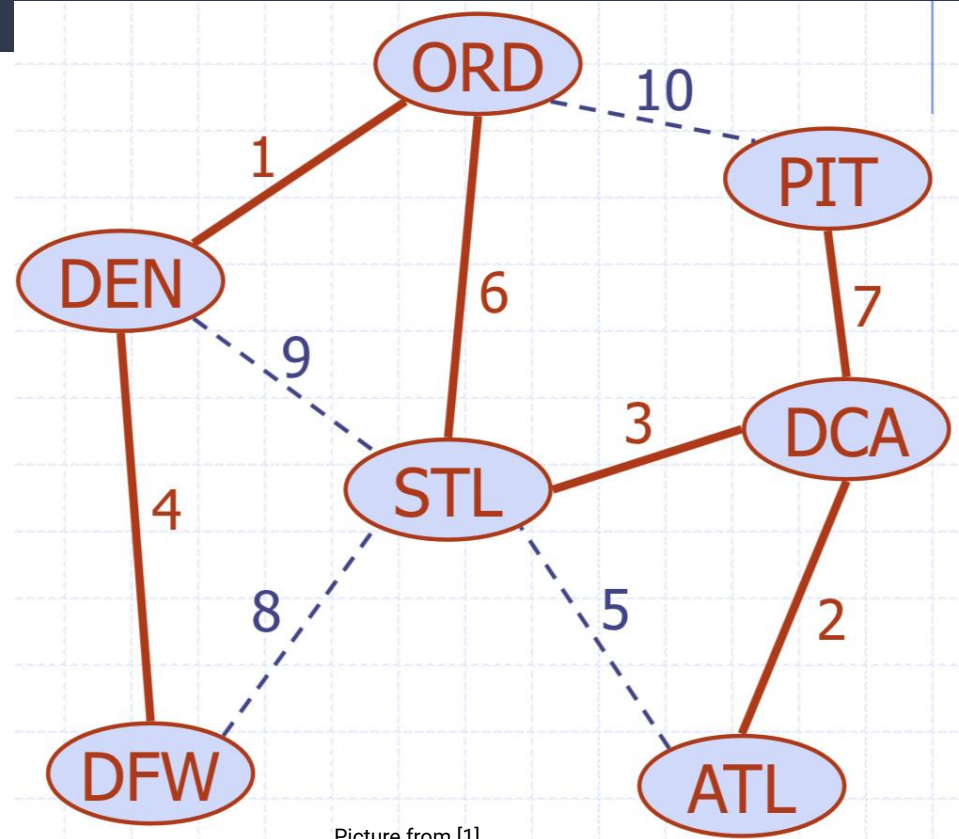
How does this problem differ from a shortest path problem?

In graph terms, how would you describe what we are looking for?

What we really want is a minimum spanning tree (MST)...

A minimum spanning tree of a graph G...

- ...is a spanning subgraph (i.e. it contains all the vertices of G)
- ...is a tree (i.e. it has no cycles)
- ...is minimal weight-wise (i.e. the total weight of all the edges it uses is minimized so that no other spanning tree has a smaller total weight)
- ...does NOT guarantee the shortest path between two vertices!




Picture from [1]

Property 1: The Cycle Property

- Let T be a minimum spanning tree of a weighted graph G
- Let e be an edge of G that is not in T
- Let C be the cycle formed by adding e to T
- Claim: For every edge f of C :
 $weight(f) \leq weight(e)$


Proof by Contradiction: ???



Property 1: The Cycle Property

- Let T be a minimum spanning tree of a weighted graph G
- Let e be an edge of G that is not in T
- Let C be the cycle formed by adding e to T
- Claim: For every edge f of C :
 $\text{weight}(f) \leq \text{weight}(e)$


Proof by Contradiction: Given all the properties above, assume that for some edge f in C , $\text{weight}(f) > \text{weight}(e)$. Then replacing f with e will produce a spanning tree T' such that the total weight of T' is smaller than the total weight of T . But that contradicts the definition of T as an MST of G .



Property 1: The Partition Property

- Consider a partition of a graph G into subsets U and V
- Let e be an edge of minimum weight across the partition (i.e. e has one endpoint in U and one endpoint in V)
- Claim: There is a minimum spanning tree of G that includes edge e


Proof: ???




Property 1: The Partition Property

- Consider a partition of a graph G into subsets U and V
- Let e be an edge of minimum weight across the partition (i.e. e has one endpoint in U and one endpoint in V)
- Claim: There is a minimum spanning tree of G that includes edge e


Proof: Let T be an MST of G and let the definitions above be true. If T does not contain e , consider the cycle C formed by e with T and let f be an edge of C across the partition. By the cycle property, $\text{weight}(f) \leq \text{weight}(e)$. Thus, $\text{weight}(f) = \text{weight}(e)$, and we obtain another MST by replacing f with e .



So how would you find an MST of this graph?



graph




MST



A 250-vertex Euclidean graph (with 1,273 edges) and its MST


Prim's Algorithm (aka Jarnik's Algorithm)

- Grow the tree one edge at a time
- Start with a single vertex (counts as a tree)
- Add $|V| - 1$ edges to it by adding the minimum-weight edge that connects a new vertex to the tree (crossing the partition that is defined by the current tree vertices--called a *crossing edge*)
- NOTE: When you add an edge to the tree, you are also adding a vertex to the tree.
- Like Dijkstra's Shortest Path, this is a greedy algorithm.



Prim's Algorithm (aka Jarnik's Algorithm)

Proof of Correctness: Follows directly from the Partition Property because we are choosing the minimum weight edge across the tree-defined partition.



Picture from [2]

Implementation of Prim's Algorithm

- How would you implement this algorithm?
- What data structures would be useful?
- What would be the space and runtime requirements?
- What would be the best way to represent the graph?

Implementation of Prim's Algorithm

- **marked[]**: an array of booleans to keep track of vertices on the tree
- **edgeTo[]**: an array to keep track of the lightest edge connecting a new vertex to the tree
- **distTo[]**: an array to keep track of the weight of the lightest edge connecting a new vertex to the tree
- **pq**: a minimum priority queue to keep track of eligible crossing edges (key is the weight of the edges)

Algorithm *PrimsMST*(G)

Input: $G = (V, E)$, a weighted, undirected graph

Output: A minimum-weight spanning tree of G

$edgeTo[]$:= a $|V|$ -sized array to store the edge connecting a vertex to the tree

$distTo[]$:= a $|V|$ -sized array to keep track of the distance of the edge connecting a vertex to a tree

$marked[]$:= a $|V|$ -sized array to keep track of which vertices have been visited

pq := a heap-based min priority queue with weights as keys and vertices as values

//initialize structures

for all $v \in V$:

$distTo[v] = \infty$

$distTo[0] = 0$

$pq.insert(0, 0)$

//main loop

while ! $pq.isEmpty$:

$v = pq.delMin()$

$marked[v] = T$

for each edge $e = (v, u)$ adjacent to v :

if $marked[u]$

continue

if $e.weight() < distTo[u]$:

$edgeTo[u] = e$


$distTo[u] = e.weight()$

if $pq.contains(u)$:

$pq.changeKey(u, distTo[u])$

else:

$pq.insert(u, distTo[u])$



//main loop

while !pq.isEmpty:

$v = pq.delMin()$

$marked[v] = T$

for each edge $e = (v, u)$ adjacent to v :

if $marked[u]$

continue

if $e.weight() < distTo[u]$:

$edgeTo[u] = e$

$distTo[u] = e.weight()$

if $pq.contains(u)$:

$pq.changeKey(u, distTo[u])$

else:

$pq.insert(u, distTo[u])$

Algorithm *PrimsMST*(G)

Input: $G = (V, E)$, a weighted, undirected graph

Output: A minimum-weight spanning tree of G

$edgeTo[]$:= a $|V|$ -sized array to store the edge connecting a vertex to the tree

$distTo[]$:= a $|V|$ -sized array to keep track of the distance of the edge connecting a vertex to a tree

$marked[]$:= a $|V|$ -sized array to keep track of which vertices have been visited

pq := a heap-based min priority queue with weights as keys and vertices as values

//initialize structures

for all $v \in V$:

$distTo[v] = \infty$

$distTo[0] = 0$

$pq.insert(0, 0)$

$O(|V|)$

//main loop

while ! $pq.isEmpty$:

$v = pq.delMin()$

$marked[v] = T$

for each edge $e = (v, u)$ adjacent to v :

if $marked[u]$

continue

if $e.weight() < distTo[u]$:

$edgeTo[u] = e$

$distTo[u] = e.weight()$

if $pq.contains(u)$:

$pq.changeKey(u, distTo[u])$

else:

$pq.insert(u, distTo[u])$

```
//main loop
```

```
while !pq.isEmpty:
```

```
    v = pq.delMin()
```

```
    marked[v] = T
```

```
    for each edge e = (v, u) adjacent to v:
```

```
        if marked[u]
```

```
            continue
```

```
        if e.weight() < distTo[u]:
```

```
            edgeTo[u] = e
```

```
            distTo[u] = e.weight()
```

```
            if pq.contains(u):
```

```
                pq.changeKey(u, distTo[u])
```

```
            else:
```

```
                pq.insert(u, distTo[u])
```

Assuming adjacency list representation, total time is **$O((V + E)\log V) = O(E\log V)$** for a connected graph

- delMin in a heap-based PQ: $O(\log V)$
- checking if PQ contains a vertex and changing a key is $O(\log V)$ as long as there is an auxiliary data structure keeping track of positions in the queue
- key of any vertex v is updated at most $\deg(v)$ times and sum of all degrees is $O(M)$

Imagine...

You are employed at a company that has offices in several different cities. Your first task is to work with a phone company to ensure that there is a connection between every pair of offices according to the following guidelines:

- The connection between two offices does not have to be direct.
- The company wants to minimize the total cost of all the connections.

What kind of a problem is this?

Imagine...


You are employed at a company that has offices in several different cities. Your first task is to work with a phone company to ensure that there is a connection between every pair of offices according to the following specifications:

- The phone company charges different rates to connect pairs of cities, specified on the following slide.
- The connection between two offices does not have to be direct.
- The company wants to minimize the total cost of all the connections.

What kind of a problem is this?

A minimum spanning tree problem!


	London	Edinburgh	Miami	Mexico City	Los Angeles
London	0	\$120	\$450	\$1022	\$830
Edinburgh	\$120	0	\$500	\$1200	\$965
Miami	\$450	\$500	0	\$210	\$197
Mexico City	\$1022	\$1200	\$210	0	\$201
Los Angeles	\$830	\$965	\$197	\$201	0



Goal: Find an MST of the graph using Kruskal's Algorithm...

Kruskal's Algorithm

- Process edges, adding edges to the tree smallest-weight first **as long as the new edge does not form a cycle.**
- The result is that the algorithm creates a forest that eventually merges into a single tree.
- What would be some of the challenges in implementing this?
- What data structures would be useful?
- How would you represent the graph?




Trace of Kruskal's algorithm

Kruskal's Algorithm

- Can be implemented using disjoint sets for each separate component
- Start with $|V|$ disjoint sets, each containing a single vertex
- Combine sets (union) by adding edges, reducing the number of separate components until there is only one

https://drive.google.com/file/d/1EJpFkGbcih3gT9nS1dVfZFU26_0lhf3/view?usp=sharing



Algorithm *KruskalsMST*(G)

Input: $G = (V, E)$, a connected, weighted, undirected graph

Output: A minimum-weight spanning tree of G

$pq :=$ a minimum priority queue of edges where keys are weights

for each edge $e = (u, v) \in E$:

$pq.insert(e, e.weight())$

$A = \emptyset$

for each $v \in V$:

$makeSet(v)$

while $|A| < |V| - 1$

$e = (u, v) = pq.delMin()$

if $findSet(u) \neq findSet(v)$:

$A = A \cup (u, v)$

$union(u, v)$

return A

Algorithm *KruskalsMST*(G)

Input: $G = (V, E)$, a connected, weighted, undirected graph

Output: A minimum-weight spanning tree of G

$pq :=$ a minimum priority queue of edges where keys are weights

for each edge $e = (u, v) \in E$:

$pq.insert(e, e.weight())$

$A = \emptyset$

for each $v \in V$:

$makeSet(v)$

while $|A| < |V| - 1$

$e = (u, v) = pq.delMin()$

if $findSet(u) \neq findSet(v)$:

$A = A \cup (u, v)$

$union(u, v)$

return A

- **$O(E \log E)$** for doing **E** inserts into the **E** -sized PQ
- Alternatively, we could just sort the edges in **$O(E \log E)$** time

Algorithm *KruskalsMST*(G)

Input: $G = (V, E)$, a connected, weighted, undirected graph

Output: A minimum-weight spanning tree of G

$pq :=$ a minimum priority queue of edges where keys are weights

for each edge $e = (u, v) \in E$:

$pq.insert(e, e.weight())$

$A = \emptyset$

for each $v \in V$:

$makeSet(v)$

while $|A| < |V| - 1$

$e = (u, v) = pq.delMin()$

if $findSet(u) \neq findSet(v)$:

$A = A \cup (u, v)$



$union(u, v)$

return A

- $makeSet$ is an $O(1)$ operation, so the total time is **$O(V)$**
- The total time required for **E union** and $findSet$ operations is **$O(E \log E)$**
- **Total time: $O(E \log E)$**

Prim's vs. Kruskal's

- Space: $|V|$ vs. $|E|$
- Time: $|E|\log|V|$ vs. $|E|\log|E|$



Baruvka's Algorithm

- ◆ Like Kruskal's Algorithm, Baruvka's algorithm grows many "clouds" at once.

Algorithm *BaruvkaMST*(G)

$T \leftarrow V$ {just the vertices of G }

while T has fewer than $n-1$ edges **do**

for each connected component C in T **do**

 Let edge e be the smallest-weight edge from C to another component in T .

if e is not already in T **then**

 Add edge e to T

return T

- ◆ Each iteration of the while-loop halves the number of connected components in T .
 - The running time is $O(m \log n)$.


Imagine...

You are employed at a florist shop as a driver during their busy Valentine's Day season. Your first task is to take a list of delivery addresses and plan a route according to the following specifications:

- You want to start and end at the florist shop.
- You want to deliver ALL the flowers
- You want to minimize the distance you have to go.
- You don't want to visit any stop more than once.
- You don't want to drive the same road more than once.


How would you go about solving this problem?


	Florist	A	B	C	D	E
Florist	0	5.5	7.3	2.1	0.9	5.8
A	5.5	0	4.2	6.5	0.12	8.2
B	7.3	4.2	0	1.2	4.3	0.5
C	2.1	6.5	1.2	0	3.2	4.4
D	0.9	0.12	4.3	3.2	0	2.5
E	5.8	8.2	0.5	4.4	2.5	0



The Traveling Salesperson Problem


- A **tour** of a graph is a spanning cycle (goes through all the vertices)
- A TSP tour of a weighted graph is a tour that is simple (i.e. repeats no vertices or edges) and has minimum weight
- Determining if a TSP tour shorter than a given length L **exists** in a graph is NP-Complete (i.e. solution can be verified in polynomial time but not discovered in polynomial time—at least not yet)
- Finding a TSP tour is NP-hard





$P \neq NP$

Complexity



$P = NP$

Other NP-Complete Problems

- Boolean Satisfiability Problem (SAT)
- Knapsack Problem
- Hamiltonian Path Problem
- Subgraph Isomorphism Problem
- Subset sum problem
- Clique Problem
- Vertex Cover Problem
- Independent set problem
- Dominating set problem
- Graph Coloring problem

References

- [1] Goodrich and Tamassia
- [2] Sedgewick and Wayne
- [3] en.wikipedia.org